

# **Distributed Storage Systems**

Compression & Delta Encoding

# Agenda



## Today's topics

- Basics of compression
- Compression schemes
- Delta encoding

# Class Structure

	Lecture	Lab
Week 1	Course introduction, networking basics, socket programming	Python sockets
Week 2	RPC, NFS, Practical RPC	Flask, JsonRPC, REST API
Week 3	AFS, reliable storage introduction	ZeroMQ, ProtoBuf
Week 4	Hard drives, RAID levels	RPi stack intro, RPi RAID with ZMQ
Week 5	Finite fields, Reed-Solomon Codes	Kodo intro, RS and RLNC with Kodo
Week 6	Repair problem, RS vs Regenerating codes	RPi simple distributed storage with Kodo RS
Week 7	Regenerating codes, XORBAS	RPi Regenerate lost fragments with RS
Week 8	Hadoop	RPi RLNC, recovery with recode
Week 9	Storage Virtualization, Network Attached Storage, Storage Area Networks	RPi basic HDFS (namenode+datanode, read & write pipeline)
Week 10	Object Storage	RPi basic S3 API
<b>Week 11</b>	<b>Compression, Delta Encoding</b>	<b>Mini project consultation</b>
Week 12	Data Deduplication	RPi Dedup
Week 13	Fog storage	Mini project consultation
Week 14	Security for Storage Systems and Recap	Mini project consultation


# **Basics of Compression**

# Key question: what is information?

- Example: you fight Muhammad Ali
  - He throws a jab with probability 0.5
  - He throws a cross with probability 0.4
  - He throws an uppercut and hook with probability 0.05 each

What carries more information if you need to block/evade?

How many bits to identify each case efficiently?



# Key question: what is information?

- Any quantification of “information” about an event depends on its probability
  - The greater the probability of an event, the smaller the information associated with knowing that the event occurred
  - Semantics are not important for this definition (but can be for actual algorithms)

# Key question: what is information?

Information of an event given that it has a probability  $p$  of happening:

$$I = \log_2(1/p) = -\log_2(p)$$

- Base 2 because we use bits :)
- Provides nice properties:
  - Information of independent events A and B is added, but probability should be multiplied
  - $I_A + I_B = -\log_2(p_A) - \log_2(p_B) = -\log_2(p_A p_B)$

# (Shannon) Entropy

Expected information in a set of possible outcomes or mutually exclusive events:

$$E[I] = H(p_1, p_2, \dots, p_N) = \sum p_i \log_2(1/p_i)$$

Another way to think of it: expected uncertainty associated with this set of events

- H is non-negative
- $H \leq \log_2(N)$

(equality when all events are equally likely)



# **Source Encoding/Compression**

# Types

**Lossless Compression** compressed can be reconstituted (uncompressed) without loss of detail or information

Example: Files, Logs, ...

**Lossy Compression** aim is to obtain the best possible fidelity for a given bit-rate or minimizing the bit-rate to achieve a given fidelity measure

Example: Video, audio

# Our Focus: Lossless Compression

When a message is drawn from a set of possible messages, each occurring with some probability

- Entropy:
  - Expected amount of information in a message
  - Lower bound on the amount of information that must be sent


# The Shannon-Fano Algorithm

This is a basic information theoretic algorithm

<b>Symbol</b>	A	B	C	D	E
<b>Count</b>	15	7	6	6	5


# Encoding for the Shannon-Fano Algorithm:

1. Sort symbols according to their frequencies/probabilities: ABCDE
2. Recursively divide into two parts, each with approx. same number of counts



# Encoding for the Shannon-Fano Algorithm:

1. Sort symbols according to their frequencies/probabilities: ABCDE
2. Recursively divide into two parts, each with approx. same number of counts



**Bits in average:**  $2 \times (15 + 7 + 6) / (39) + 3 \times (6 + 5) / 39 \sim 2,28$  bits/symbols

# Huffman Coding

Based on the probability (or frequency of occurrence) of a data item

The principle is to use a lower number of bits to encode the data that occurs more frequently

Codes are stored in a **Code Book**

The code book + encoded data must be transmitted to enable decoding

Fixed # of bits to variable # of bits

# Huffman Coding

- **Unique Prefix Property:**
  - No code is a prefix to any other code
  - Great for decoder, unambiguous
- If prior statistics are available and accurate, then Huffman coding is very good



# Huffman Coding

1. **Init:** Put all nodes in an OPEN list, keep it sorted at all times (e.g., ABCDE)
2. Repeat until the OPEN list has only one node left:
  - a. From OPEN pick two nodes having the lowest frequencies/probabilities, create a parent node of them
  - b. Assign the sum of the children's frequencies/probabilities to the parent node and insert it into OPEN
  - c. Assign code 0, 1 to the two branches of the tree, and delete the children from OPEN

# Huffman Coding

Count	Symbol
-------	--------


15	A
----	---

7	B
---	---

6	C
---	---

6	D
---	---

5	E
---	---



# Huffman Coding

Count	Symbol
-------	--------

15	A
----	---

7	B
---	---


6	C
---	---

6	D
---	---

5	E
---	---

13

11



# Huffman Coding

**Count**   **Symbol**

15

A

15

7

B

13

6

C

24


6

D

11

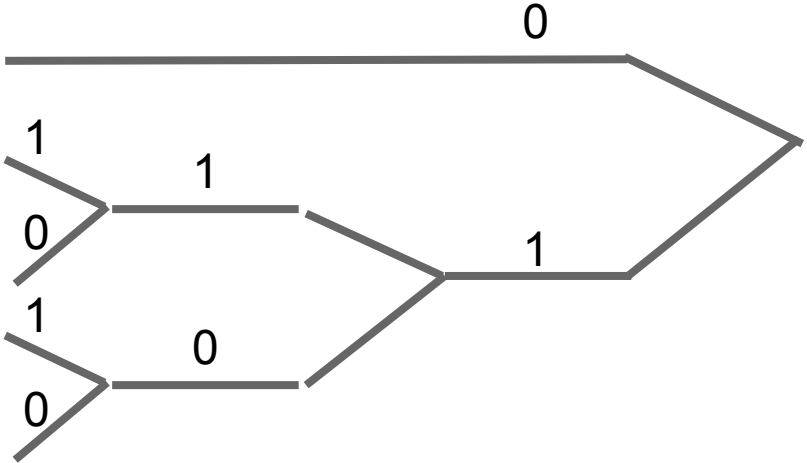
5

E



# Huffman Coding

Count	Symbol		Code
15	A	0	0
7	B	1	111
6	C	0	110
6	D	1	101
5	E	0	100



The diagram illustrates the construction of a Huffman tree. It starts with five symbols: A (count 15), B (count 7), C (count 6), D (count 6), and E (count 5). The tree is built by repeatedly merging the two smallest nodes. First, B and C are merged into a node with count 13, with B as the left child (labeled 1) and C as the right child (labeled 0). Then, D and E are merged into a node with count 11, with D as the left child (labeled 1) and E as the right child (labeled 0). Finally, the node containing B and C (count 13) and the node containing D and E (count 11) are merged into the root node with count 24. The root node has a left child (labeled 0) which is symbol A, and a right child (labeled 1) which is the internal node containing the subtree for B, C, D, and E. The subtree for B, C, D, and E has a left child (labeled 1) which is the internal node containing B and C, and a right child (labeled 0) which is the internal node containing D and E. The subtree for B and C has a left child (labeled 1) which is symbol B, and a right child (labeled 0) which is symbol C. The subtree for D and E has a left child (labeled 1) which is symbol D, and a right child (labeled 0) which is symbol E.

# Huffman Coding

**Count    Symbol**

**15**

**(Jab)**

7

6

6

5


**A**

**B**

**C**

**D**

**E**




**Code**

**0**


111

110

101




# Huffman Coding



**Bits in average:**  $1 \times 15/(39) + 3 \times (24)/39 \sim 2,2307$  bits/symbol

How about entropy?

# Huffman Coding




**Bits in average:**  $1 \times 15/(39) + 3 \times (24)/39 \sim 2,2307$  bits/symbol

**Entropy:**

$$\begin{aligned}
 & - 15/39 \times \log_2(15/39) - 7/39 \times \log_2(7/39) - 2 \times 6/39 \times \log_2(6/39) - 5/39 \times \log_2(5/39) \\
 & = 0,5301967 + 0,44477772 + 0,8309045 + 0,3799325 \\
 & = \mathbf{2,18581142}
 \end{aligned}$$



# Huffman Coding



**Bits in average:**  $1 \times 15/(39) + 3 \times (24)/39 \sim 2,2307$  bits/symbol

**Entropy:**

$$\begin{aligned}
 & - 15/39 \times \log_2(15/39) - 7/39 \times \log_2(7/39) - 2 \times 6/39 \times \log_2(6/39) - 5/39 \times \log_2(5/39) \\
 & = 0,5301967 + 0,44477772 + 0,8309045 + 0,3799325 \\
 & = \mathbf{2,18581142}
 \end{aligned}$$

**Shannon-Fano:**  $\sim 2,28$  bits/symbols

# Adaptive Huffman Coding

The basic Huffman algorithm has been extended

- Statistical knowledge is often not available
- Even when it is available, it could be a heavy overhead especially when many tables have to be sent

The solution is to use adaptive algorithms, e.g. Adaptive Huffman coding (applicable to other adaptive compression algorithms).

## **Problem:** how to determine probabilities?

- Simple idea is to use adaptive model:
  - Start with guess of symbol frequencies
  - Update frequency with each new symbol
- Another idea is to take account of intersymbol probabilities, e.g., prediction by partial matching

# Lempel-Ziv-Welch (LZW) Algorithm

The LZW algorithm is a very common compression technique

Example: Suppose we want to encode the Oxford Concise English dictionary which contains about 159,000 entries

Why not just transmit each word as an 18 bit number?

# Lempel-Ziv-Welch (LZW) Algorithm

The LZW algorithm is a very common compression technique

Example: Suppose we want to encode the Oxford Concise English dictionary which contains about 159,000 entries

Why not just transmit each word as an 18 bit number?

- Too many bits
- Everyone needs a dictionary
- For this simple example, it would work for English text only

# Lempel-Ziv-Welch (LZW) Algorithm

The LZW algorithm is a very common compression technique


Example: Suppose we want to encode the Oxford Concise English dictionary which contains about 159,000 entries

Why not just transmit each word as an 18 bit number?

- **Solution:** Find a way to build the dictionary adaptively
  - Original methods due to Ziv and Lempel in 1977 and 1978. Terry Welch improved the scheme in 1984 (called LZW compression).
  - Variable length to fixed length (“reverse” of Huffman)

# LZW Compression

- Table size:  $2^N$
- N bits to represent each index
- Send **table indexes**
- String table can be reconstructed by the decoder using information from the encoded stream → table is not sent



# LZW Compression Algorithm

```
w = get input symbol;  
while ( there are still input symbols )  
{  
    k = read a character  
    if w + k exists in the dictionary  
        w = w + k;  
    else  
        add w+k to the dictionary;  
        output the code for w;  
        w = k;  
}
```

Output the code for w

(+ = concatenate)




# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;
while ( there are still input symbols )
{
  k = read a character
  if w + k exists in the dictionary
    w = w + k;
  else
    add w+k to the dictionary;
    output the code for w;
    w = k;
}
Output the code for w
```

Output:

**w = a**



# LZW Compression Algorithm


***abbbabbbab***

```
w = get input symbol;  
while ( there are still input symbols )  
{  
  k = read a character  
  if w + k exists in the dictionary  
    w = w + k;  
  else  
    add w+k to the dictionary;  
    output the code for w;  
    w = k;  
}  
Output the code for w
```

Output:

w = a

k = b



# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;  
while ( there are still input symbols )  
{  
    k = read a character  
    if w + k exists in the dictionary  
        w = w + k;  
    else  
        add w+k to the dictionary;  
        output the code for w;  
        w = k;  
}  
Output the code for w
```

Output: **97**,

w = a


k = b

add w+k to the dictionary; **ab**

**output the code for w;**

**w = k;**

**w=b**




# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;  
while ( there are still input symbols )      w = b  
  {  
    k = read a character  
    if w + k exists in the dictionary  
      w = w + k;  
    else  
      add w+k to the dictionary;  
      output the code for w;  
      w = k;  
  }  
Output the code for w
```

Output: 97,



# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;
while ( there are still input symbols )
{
    k = read a character
    if w + k exists in the dictionary
        w = w + k;
    else
        add w+k to the dictionary;
        output the code for w;
        w = k;
}
Output the code for w
```

Output: **97**,

w = b

k = b

	⋮
97	a
98	b
256	ab
257	
258	
	⋮
$2^N-3$	
$2^N-2$	
$2^N-1$	

# LZW Compression Algorithm

*abbbabbbbab*

```

w = get input symbol;
while ( there are still input symbols )
{
    k = read a character
    if w + k exists in the dictionary
        w = w + k;
    else
        add w+k to the dictionary;
        output the code for w;
        w = k;
}
Output the code for w
    
```


Output: **97, 98,**

w = b

k = b

bb

w=b




# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;
while ( there are still input symbols )      w = b
{
  k = read a character
  if w + k exists in the dictionary
    w = w + k;
  else
    add w+k to the dictionary;
    output the code for w;
    w = k;
}
Output the code for w
```

Output: 97, 98,



# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;  
while ( there are still input symbols )  
{  
  k = read a character  
  if w + k exists in the dictionary  
    w = w + k;  
  else  
    add w+k to the dictionary;  
    output the code for w;  
    w = k;  
}  
Output the code for w
```

w = b

k = b

	⋮
97	a
98	b
256	ab
257	bb
258	
	⋮
$2^N-3$	
$2^N-2$	
$2^N-1$	

Output: 97, 98,



# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;
while ( there are still input symbols )
{
    k = read a character
    if w + k exists in the dictionary
        w = w + k;
    else
        add w+k to the dictionary;
        output the code for w;
        w = k;
}
```

Output the code for w

Output: **97**, **98**,

w = b

k = b

**w=bb**

	⋮
97	a
98	b
256	ab
257	bb
258	
	⋮
$2^N-3$	
$2^N-2$	
$2^N-1$	

# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;  
while ( there are still input symbols )  
{  
  k = read a character  
  if w + k exists in the dictionary  
    w = w + k;  
  else  
    add w+k to the dictionary;  
    output the code for w;  
    w = k;  
}  
Output the code for w
```

**w = bb**

**k = a**

	⋮
97	a
98	b
256	ab
257	bb
258	
	⋮
$2^N-3$	
$2^N-2$	
$2^N-1$	

Output: **97**, **98**,

# LZW Compression Algorithm

*abbbabbbbab*

```
w = get input symbol;  
while ( there are still input symbols )  
{  
  k = read a character  
  if w + k exists in the dictionary  
    w = w + k;  
  else  
    add w+k to the dictionary;  
    output the code for w;  
    w = k;  
}  
Output the code for w
```

Output: **97, 98, 257,**

w = bb

k = a

bba

w=a

	⋮
97	a
98	b
	⋮
256	ab
257	bb
258	bba
	⋮
$2^N-3$	
$2^N-2$	
$2^N-1$	

# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;  
while ( there are still input symbols )  
{  
  k = read a character  
  if w + k exists in the dictionary  
    w = w + k;  
  else  
    add w+k to the dictionary;  
    output the code for w;  
    w = k;  
}  
Output the code for w
```

**w = a**

**k = b**

	⋮
97	a
98	b
256	ab
257	bb
258	bba
	⋮
$2^N-3$	
$2^N-2$	
$2^N-1$	

Output: **97, 98, 257**

# LZW Compression Algorithm

*abbbabbbbab*

```
w = get input symbol;  
while ( there are still input symbols )  
{  
    k = read a character  
    if w + k exists in the dictionary  
        w = w + k;  
    else  
        add w+k to the dictionary;  
        output the code for w;  
        w = k;  
}  
Output the code for w
```

Output: **97**, **98**, **257**,

w = a

k = b

**w=ab**

	⋮
97	a
98	b
	⋮
256	ab
257	bb
258	bba
	⋮
$2^N-3$	
$2^N-2$	
$2^N-1$	

# LZW Compression Algorithm

*abbbabbbbab*

```
w = get input symbol;  
while ( there are still input symbols )  
{  
  k = read a character  
  if w + k exists in the dictionary  
    w = w + k;  
  else  
    add w+k to the dictionary;  
    output the code for w;  
    w = k;  
}  
Output the code for w
```

**w = ab**

**k = b**

	⋮
97	a
98	b
256	ab
257	bb
258	bba
	⋮
$2^N-3$	
$2^N-2$	
$2^N-1$	

Output: **97, 98, 257**

# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;  
while ( there are still input symbols )  
{  
  k = read a character  
  if w + k exists in the dictionary  
    w = w + k;  
  else  
    add w+k to the dictionary;  
    output the code for w;  
    w = k;  
}  
Output the code for w
```

Output: **97, 98, 257, 256**

w = ab

k = b

abb

w=b

	⋮
97	a
98	b
	⋮
256	ab
257	bb
258	bba
259	abb
	⋮
$2^N-2$	
$2^N-1$	

# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;  
while ( there are still input symbols )  
{  
  k = read a character  
  if w + k exists in the dictionary  
    w = w + k;  
  else  
    add w+k to the dictionary;  
    output the code for w;  
    w = k;  
}  
Output the code for w
```

**w = b**

**k = b**

	⋮
97	a
98	b
	⋮
256	ab
257	bb
258	bba
259	abb
	⋮
$2^N-2$	
$2^N-1$	

Output: **97, 98, 257, 256**



# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;
while ( there are still input symbols )
{
    k = read a character
    if w + k exists in the dictionary
        w = w + k;
    else
        add w+k to the dictionary;
        output the code for w;
        w = k;
}
Output the code for w
```

w = b

k = b

**w=bb**

	⋮
97	a
98	b
256	ab
257	bb
258	bba
259	abb
	⋮
$2^N-2$	
$2^N-1$	

Output: **97, 98, 257, 256**

# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;
while ( there are still input symbols )
{
    k = read a character
    if w + k exists in the dictionary
        w = w + k;
    else
        add w+k to the dictionary;
        output the code for w;
        w = k;
}
Output the code for w
```

**w = bb**

**k = a**

	⋮
97	a
98	b
256	ab
257	bb
258	bba
259	abb
	⋮
$2^N-2$	
$2^N-1$	

Output: **97, 98, 257, 256**

# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;
while ( there are still input symbols )
{
    k = read a character
    if w + k exists in the dictionary
        w = w + k;
    else
        add w+k to the dictionary;
        output the code for w;
        w = k;
}
Output the code for w
```

w = bb

k = a

**w=bba**

	⋮
97	a
98	b
256	ab
257	bb
258	bba
259	abb
	⋮
$2^N-2$	
$2^N-1$	

Output: **97, 98, 257, 256**

# LZW Compression Algorithm

*abbabbab*

```
w = get input symbol;  
while ( there are still input symbols )  
{  
  k = read a character  
  if w + k exists in the dictionary  
    w = w + k;  
  else  
    add w+k to the dictionary;  
    output the code for w;  
    w = k;  
}  
Output the code for w
```

**w = bba**

**k = b**

	⋮
97	a
98	b
256	ab
257	bb
258	bba
259	abb
	⋮
$2^N-2$	
$2^N-1$	

Output: **97, 98, 257, 256**

# LZW Compression Algorithm

*abbbabbbab*

```

w = get input symbol;
while ( there are still input symbols )
{
    k = read a character
    if w + k exists in the dictionary
        w = w + k;
    else
        add w+k to the dictionary;
        output the code for w;
        w = k;
}

```

w = bba

k = b

add w+k to the dictionary; **bbab**

**output the code for w;**

**w = k;**

**w=b**

**Output the code for w**

Output: **97**, **98**, **257**, **256**, **258**, **98**

	⋮
97	a
98	b
	⋮
256	ab
257	bb
258	bba
259	abb
260	bbab
	⋮
$2^N-1$	

# LZW Compression Algorithm

*abbbabbbab*

10 symbols of 8 bits each → 80 bit sequence

```
w = get input symbol;
while ( there are still input symbols )
{
    k = read a character
    if w + k exists in the dictionary
        w = w + k;
    else
        add w+k to the dictionary;
        output the code for w;
        w = k;
}
```

Output: 97, 98, 257, 256, 258, 98

6 symbols of N bits each → 6N bit sequence

# LZW Compression Algorithm

*abbbabbbab*

10 symbols of 8 bits each → 80 bit sequence

```
w = get input symbol;
while ( there are still input symbols )
{
    k = read a character
    if w + k exists in the dictionary
        w = w + k;
    else
        add w+k to the dictionary;
        output the code for w;
        w = k;
}
```

As long as  $6N < 80 \rightarrow$  compression

$N < 40/3 \sim 13,333$  bits

Ex:  $N = 12$  bits  $\rightarrow$  Rate =  $6 \times 12 / 80 = 0,9$

$N_{\min} = 9$  bits  $\rightarrow$  Rate =  $6 \times 9 / 80 = 0,675$

Output: 97, 98, 257, 256, 258, 98

6 symbols of N bits each → 6N bit sequence

# LZW Compression Algorithm

*abbbabbbab*

```
w = get input symbol;
while ( there are still input symbols )
{
    k = read a character
    if w + k exists in the dictionary
        w = w + k;
    else
        add w+k to the dictionary;
        output the code for w;
        w = k;
}
```

Clearly, for this particular sequence there might have been an advantage to use a smaller starting dictionary (only containing a and b)

Best case (if you look back and try to optimize) would require 3 bits per symbol in the output: 18 bits

For such small sequence, how much of we set 1 bit for representing a or b?

Output: 97, 98, 257, 256, 258, 98



# Lossless Compression Algorithms (Pattern Substitution)

- We substitute frequently repeating patterns with a code
- The code is shorter than the pattern

A simple Pattern Substitution scheme could employ predefined code

Example: replace all occurrences of 'The' with the code '&'

# Lossless Compression Algorithms (Pattern Substitution)

More typically tokens are assigned to according to frequency of occurrence of patterns:

- Count occurrence of tokens
- Sort in Descending order
- Assign some symbols to highest count tokens

A predefined symbol table may used i.e. assign code  $i$  to token  $i$  (good if you know the application)

It is more usual to dynamically assign codes to tokens

(We will see a different example for Delta Encoding later on)

# Other Compression Techniques

## Simple Repetition Suppression

- If  $n$  successive tokens appear, we can replace them with a token and a count number of occurrences
- We usually need to have a special **flag** to denote when the repeated token appears

**Example:** 89400

can be replaced with

894f32

where **f** is the flag for zero

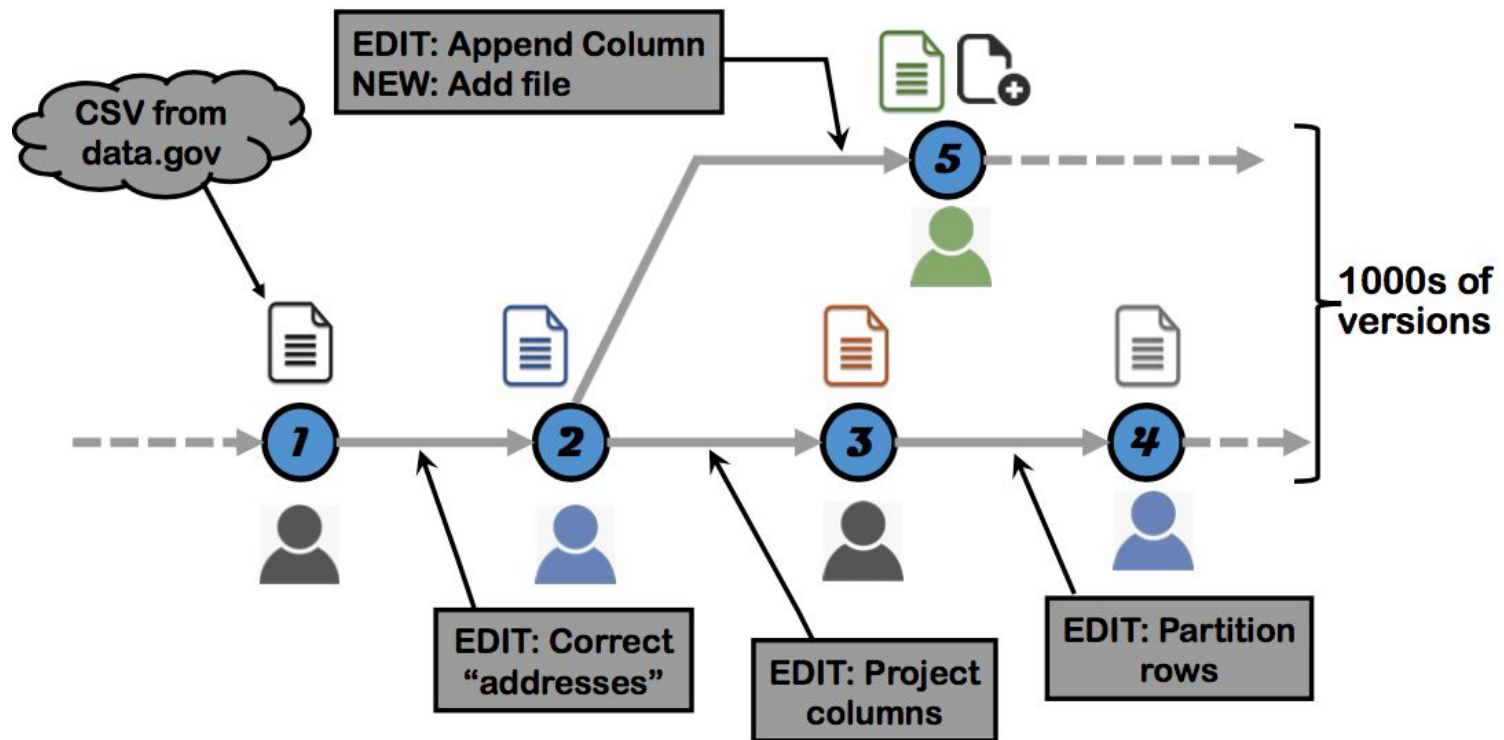
- Compression savings depend on the content of the data
- **Applications:** Zero Length Suppression for silence in audio data, pauses in conversation, blanks in text or program source files, backgrounds in images

# Run-length Encoding

- Frequently applied to images (or pixels in a scan line).
- Used in JPEG (small part of the process)
- Sequences of image elements  $X_1, X_2, \dots, X_n$  are mapped to pairs  $(c_1, l_1), (c_1, l_2), \dots, (c_n, l_n)$  where  $c_i$  represent image intensity or colour and  $l_i$  the length of the  $i$ -th run of pixels
- **Example:**  
Original Sequence: 111122233333311112222  
can be encoded as: (1,4),(2,3),(3,6),(1,4),(2,4)
- Worst case: encoding is more heavy than original file  
2\*integer rather 1\* integer if data is represented as integers

# **Data Versioning**

# A typical data analysis workflow



# The dataset versioning problem

- Many private copies of the datasets lead to **massive redundancy** in storage
- No easy way to keep **track of dependencies**
- No mechanisms to support and record manual **conflict resolution**
- No way to **analyze/compare/query** versions (across users)

# Dataset version control desiderata

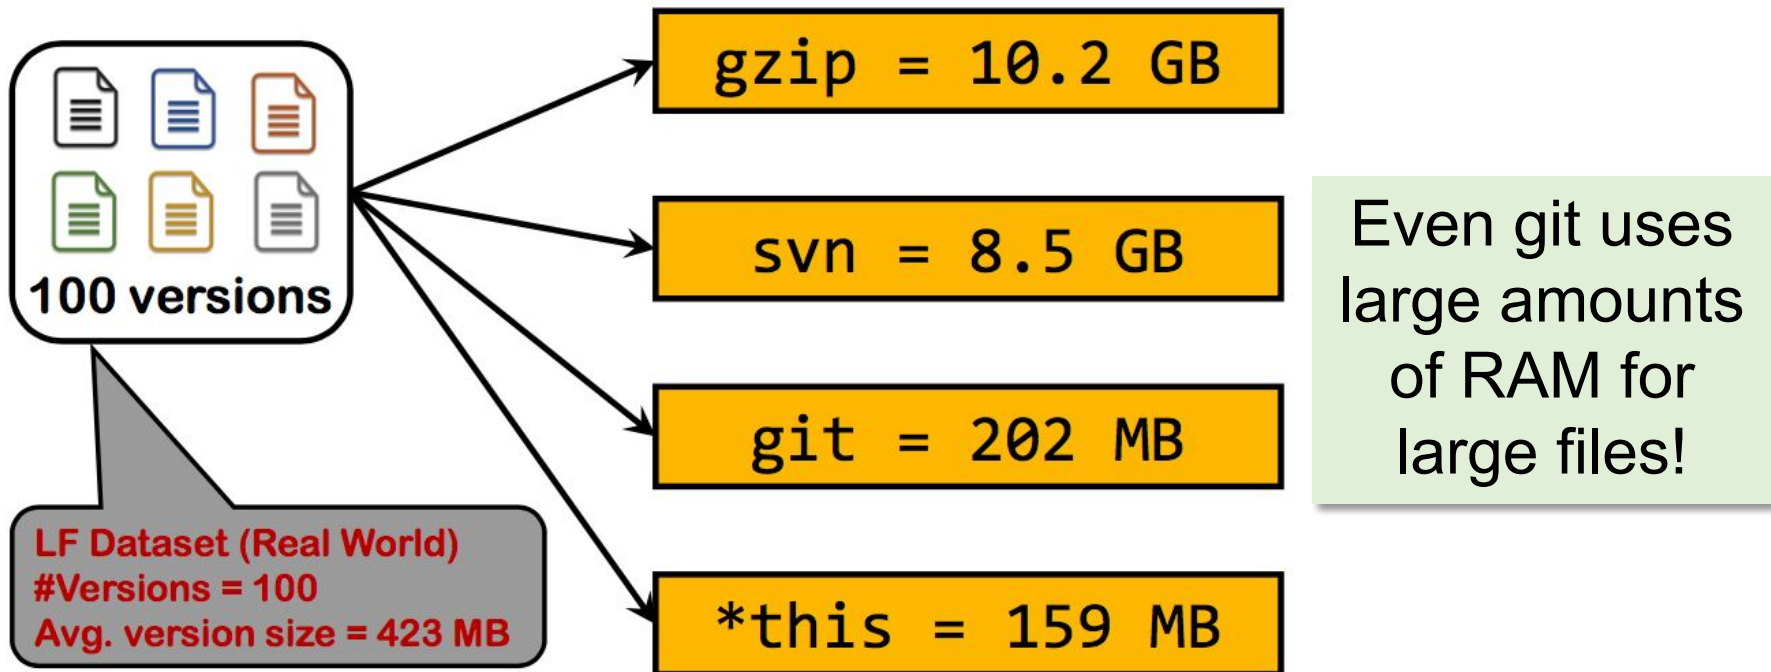
- Branch, update, merge, transform
  - Large unstructured or structured datasets
- Main challenges:
  - How can we store thousands of versions of datasets compactly?
  - How to access any version, on-demand, efficiently?



# Version Control Systems

- We already have Git/SVN and many more
- Versioning algorithms optimized to work with code-like data
  - Sparse
  - Local changes (focused in specific parts of the file)
- Scenario: What if we reformat a date that appears in all tuples of a structured dataset?

# Version Control Systems in Practice

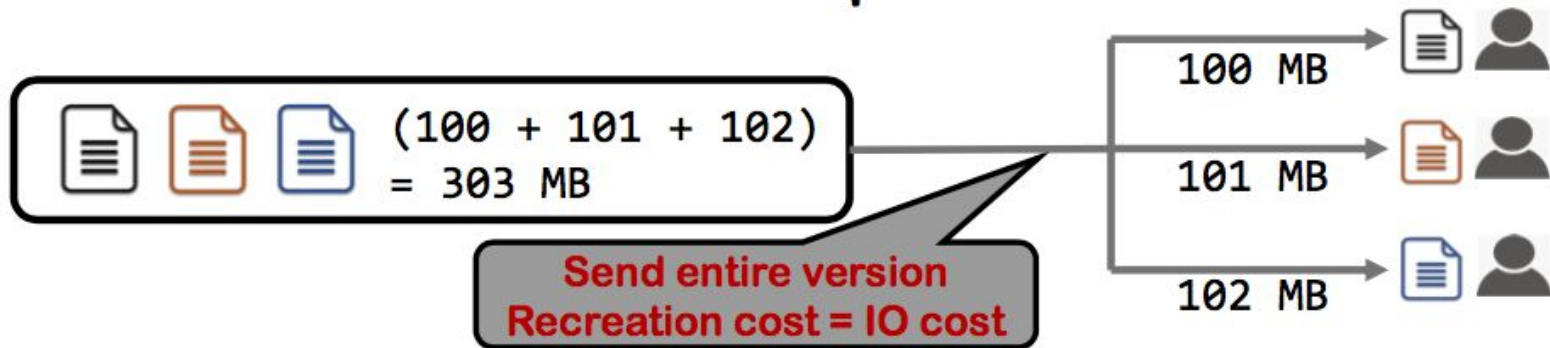


# Costs

**Storage Cost:** Space to store all versions




**Recreation Cost:** time required to access a version



# How to recreate a version?

- Use **eltas**: A delta between versions is a file which allows constructing one version given the other
- A delta has its own storage cost and recreation cost
- Examples: Unix diff, xdelta, VCDIFF



# VCDIFF ( RFC 3284 )

Uses a source (can be empty as an input) to encode an incoming file and perform delta compression

Represents a file using basic operations (including):

**COPY( $t, p$ )** defines the  $t$  number of bytes to use starting from position  $p$  in the dictionary/source

**ADD( $t, s$ )** indicates the  $t$  number of bytes to add to the dictionary (and the new file) with a sequence  $s$  ( $s$  has  $t$  bytes)

**RUN ( $x, s$ ):** A byte  $b$  will be repeated  $x$  times

# VCDIFF ( RFC 3284 )

**File  $V_1$ :** 0, 1, 2, 3, 4, 5, 6, 7, 8, 9

can be a **source** in VCDIFF terms (a prior dictionary)

**File  $V_2$ :** 0, 1, 2, 7, 8, 9, 3, 4, 5, 6, F, D

arrives, the delta representation of this file is  
{COPY (3,0), COPY(3,7), COPY(4,3), ADD(2, {F,D}) }

→ Recognized two sequences of 3 bytes & one of 4 bytes that could be used in the delta representation

# VCDIFF ( RFC 3284 )

**File  $V_1$ :** 0, 1, 2, 3, 4, 5, 6, 7, 8, 9

can be a **source** in VCDIFF terms (a prior dictionary)

**File  $V_2$ :** 0, 1, 2, 7, 8, 9, 3, 4, 5, 6, F, D

arrives, the delta representation of this file is  
{COPY (3,0), COPY(3,7), COPY(4,3), ADD(2, {F,D}) }

→ Performance is usually close to GZIP

→ Can use multiple previous files as source

# VCDIFF - Example

**Source:**

a b c d e f g h i j k l m n o p

**Target:**

a b c d w x y z e f g h e f g h e f g h e f g h z z z z



# VCDIFF - Example

Position 0

↓  
**Source:**  
a b c d e f g h i j k l m n o p

**Target:**  
a b c d w x y z e f g h e f g h e f g h z z z z

COPY 4, 0

# VCDIFF - Example

**Source:**

a b c d e f g h i j k l m n o p

**Target:**

a b c d **w x y z** e f g h e f g h e f g h e f g h z z z z

COPY 4, 0

ADD 4, w x y z

# VCDIFF - Example

Position 4

**Source:**

a b c d e f g h i j k l m n o p

**Target:**

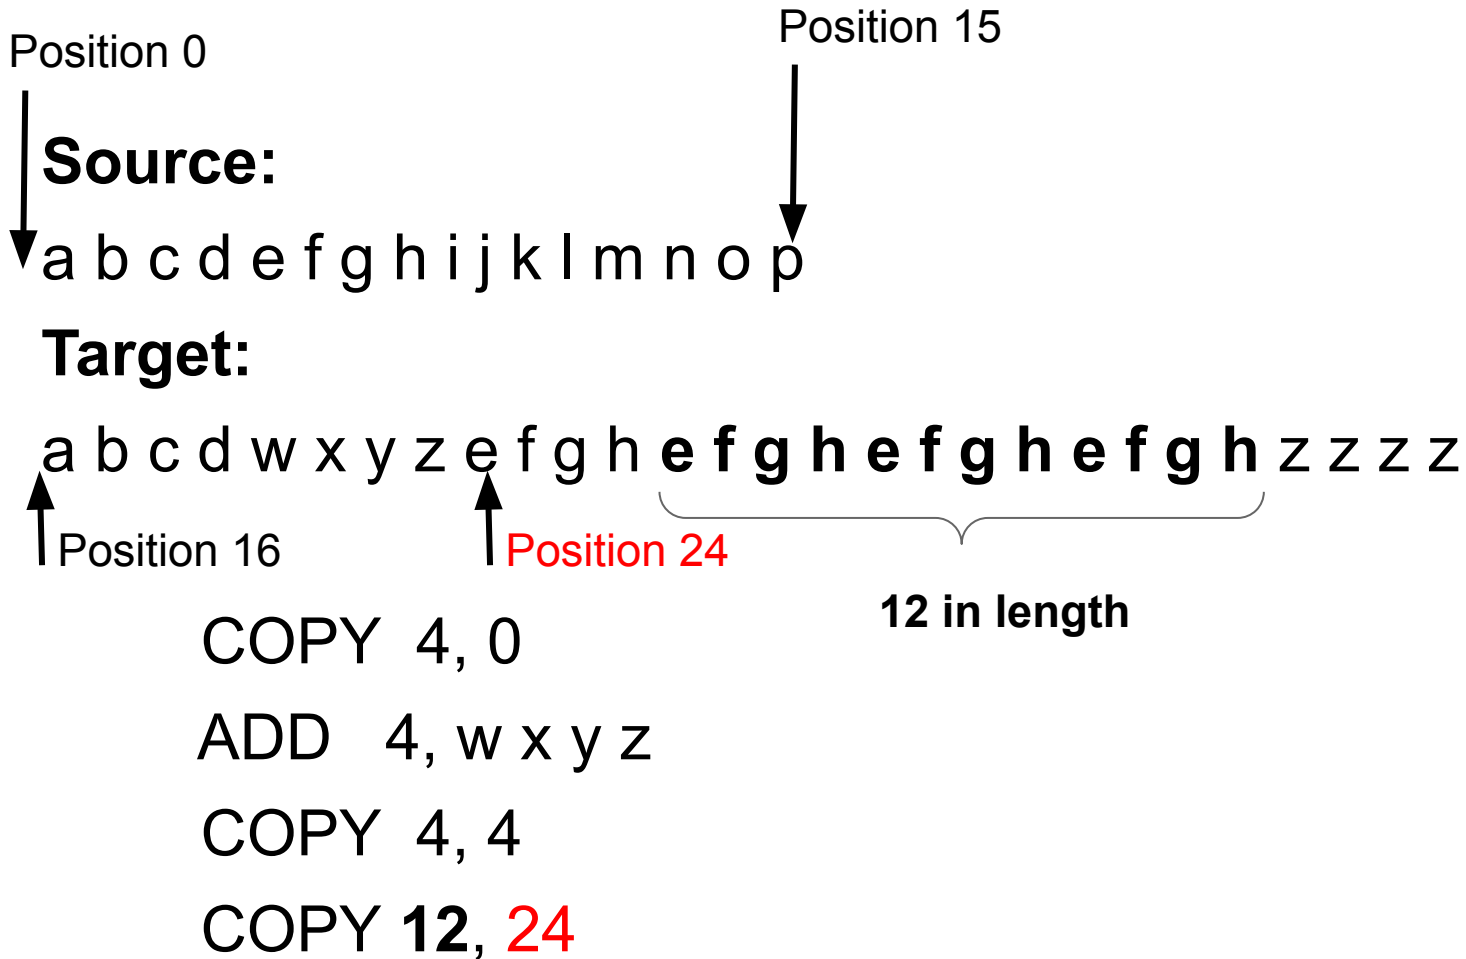
a b c d w x y z **e f g h** e f g h e f g h e f g h z z z z

COPY 4, 0

ADD 4, w x y z

COPY 4, 4

# VCDIFF - Example



# VCDIFF - Example

**Source:**

a b c d e f g h i j k l m n o p

**Target:**

a b c d w x y z e f g h e f g h e f g h e f g h **z z z z**

COPY 4, 0


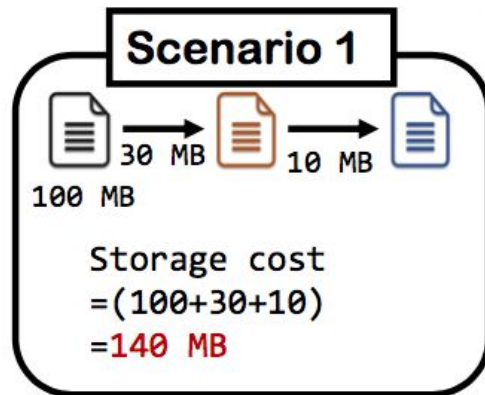
ADD 4, w x y z

COPY 4, 4


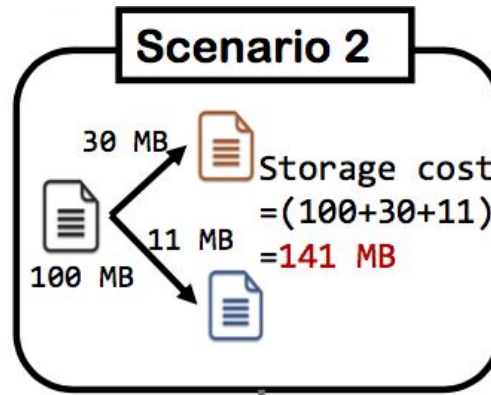
COPY 12, 24

RUN 4, z


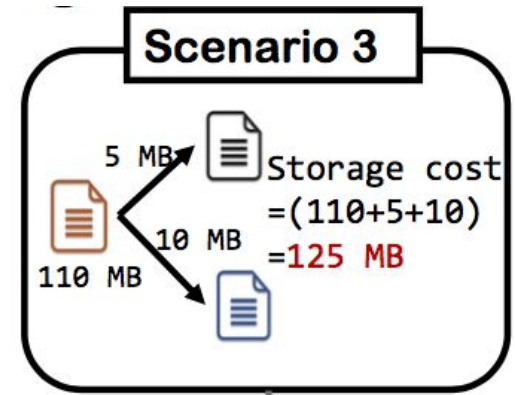
# Storage/Recreation Tradeoff (with delta encoding)



Total Access Cost  
 $= 370 \text{ MB}$



Total Access Cost  
 $= 341 \text{ MB}$



Total Access Cost  
 $= 345 \text{ MB}$

# Problem

Find a storage solution that:

- Minimizes (total) recreation cost within storage budget
- Minimizes maximum recreation cost within a storage budget

