

Distributed Storage Systems

Data Deduplication

Agenda



Today's topics

- Data deduplication
- Practical ideas
- A generalized view of data deduplication

Class Structure

	Lecture	Lab
Week 1	Course introduction, networking basics, socket programming	Python sockets
Week 2	RPC, NFS, Practical RPC	Flask, JsonRPC, REST API
Week 3	AFS, reliable storage introduction	ZeroMQ, ProtoBuf
Week 4	Hard drives, RAID levels	RPi stack intro, RPi RAID with ZMQ
Week 5	Finite fields, Reed-Solomon Codes	Kodo intro, RS and RLNC with Kodo
Week 6	Repair problem, RS vs Regenerating codes	RPi simple distributed storage with Kodo RS
Week 7	Regenerating codes, XORBAS	RPi Regenerate lost fragments with RS
Week 8	Hadoop	RPi RLNC, recovery with recode
Week 9	Storage Virtualization, Network Attached Storage, Storage Area Networks	RPi basic HDFS (namenode+datanode, read & write pipeline)
Week 10	Object Storage	RPi basic S3 API
Week 11	Compression, Delta Encoding	Mini project consultation
Week 12	Data Deduplication	RPi Dedup
Week 13	Fog storage	Mini project consultation
Week 14	Security for Storage Systems and Recap	Mini project consultation

Cost of Storing Data

- ...

Cost of Storing Data

- Individual file/data sizes
- Redundancy introduced for protection

Cost of Storing Data

- Individual file/data sizes
- Redundancy introduced for protection

Can we reduce the storage costs?

Cost of Storing Data

- Individual file/data sizes
- Redundancy introduced for protection

Can we reduce the storage costs?

- RAID / Erasure Codes / Network Codes
- ...

Cost of Storing Data

- Individual file/data sizes
- Redundancy introduced for protection

Can we reduce the storage costs?

- RAID / Erasure Codes / Network Codes
- Compress individual files
- ...

Cost of Storing Data

- Individual file/data sizes
- Redundancy introduced for protection

Can we reduce the storage costs?

- RAID / Erasure Codes / Network Codes
- Compress individual files
- Data deduplication (say what?)

Basic example

Imagine:

- You have a mail server
- An email is sent to 100 of your employees
- Email has a 10MB attachment (same for everyone)
- What is the cost of that email to you?

Basic example

Imagine:

- You have a mail server
- An email is sent to 100 of your employees
- Email has a 10MB attachment (same for everyone)
- What is the cost of that email to you?
 - ~1000MB + 100x (Rest of email) (without reliability)
 - ~3000MB + 300x (Rest of email) (3-way replication)
- Can we do better?
 - Store the 10MB

Basic example

Imagine:

- You have a mail server
- An email is sent to 100 of your employees
- Email has a 10MB attachment (same for everyone)
- What is the cost of that email to you?
 - ~1000MB + 100x (Rest of email) (without reliability)
 - ~3000MB + 300x (Rest of email) (3-way replication)
- Can we do better?
 - Store the 10MB once per receiver
 - 10MB + overhead + 100x(Rest of email)


Basic example

Imagine:

- You have a mail server
- An email is sent to 100 of your employees
- Email has a 10MB attachment (same for everyone)
- What is the cost of that email to you?
 - ~1000MB + 100x (Rest of email) (without reliability)
 - ~3000MB + 300x (Rest of email) (3-way replication)
- Can we do better?
 - Store the 10MB once per receiver
 - 10MB + overhead + 100x(Rest of email)
- Do we need to do it in a full file?



What is Deduplication?

Think of it as a LEGO problem




What is Deduplication?

Think of it as a LEGO problem

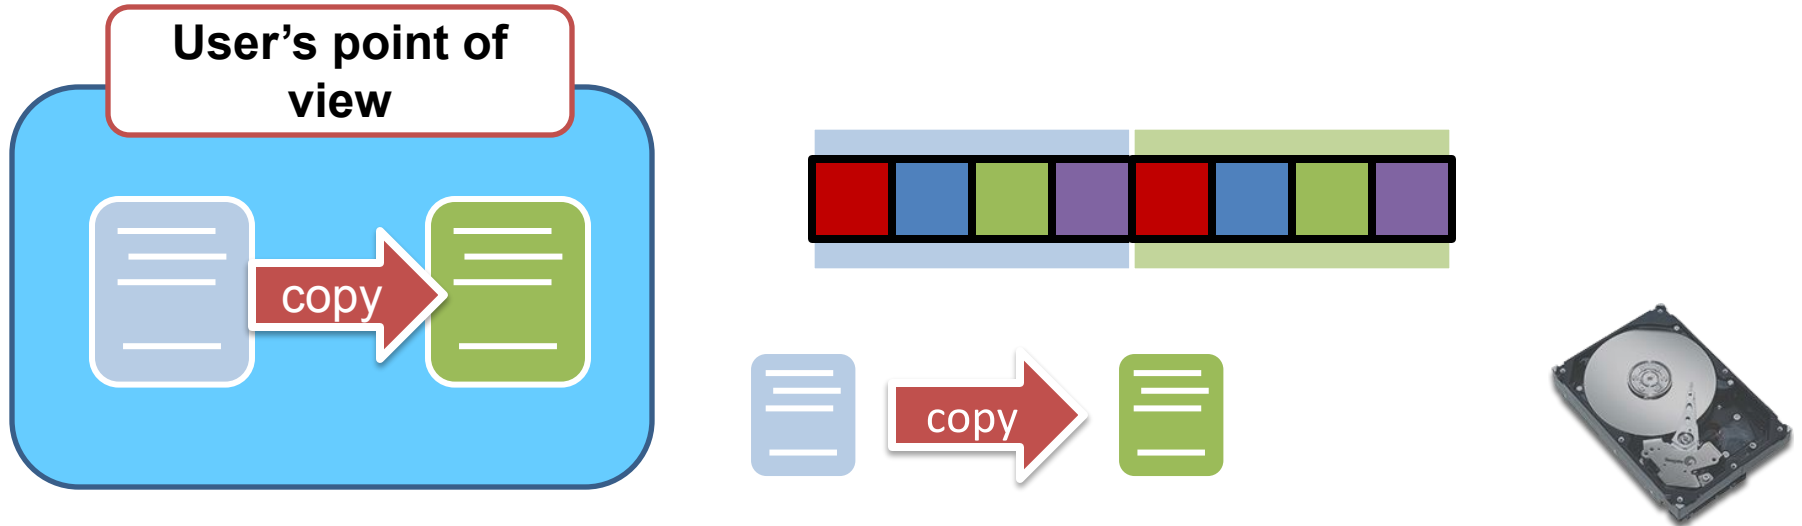


What is Deduplication?

But you can also take the data objects and convert them into pieces:





What is Deduplication?




Goal: Eliminate storage of data with same content



What is Deduplication?




Goal: Eliminate storage of data with same content




What is Deduplication?



Goal: Eliminate storage of data with same content




Block-level deduplication




Example: storage space reduced by

Block-level deduplication




Example: storage space reduced by

Block-level deduplication



Example: storage space reduced by

Block-level deduplication




Example: storage space reduced by $5/12 = 42\%$

Performance tends to improve as the system stores more data

Slide adapted from Patrick Lee


Post-process vs Inline



Process contention increases with #processes

- Copy to tape: Too slow to stream tape
- Recovery: SLA predictability
- Replication: Poor time-to-DR
- Deduplication itself if interleaved with backup or restore

More admin needed to fight these issues



Other activities unimpeded

- Predictable
- Simpler

Challenges

- Can we preserve the performance?
- Can we support general file system operations?
 - Read, write, modify, delete
- Can we deploy deduplication on low-cost commodity systems?
 - e.g., a few GB of RAM, 32/64-bit CPU, standard OS

Some Work in Deduplication

- Deduplication backup systems
 - e.g., Venti [Quinlan & Dorward '02], Data Domain [Zhu et al. '08], Foundation [Rhea et al. '08]
 - Assume data is not modified or deleted
- Deduplication file systems
 - OpenSolaris ZFS, OpenDedup SDFS
 - Consume significant memory space, not for commodity systems
- VM image storage
 - e.g., Lithium [Hansen & Jul '10], mainly on fault tolerance, but not on deduplication

Basics: Layout

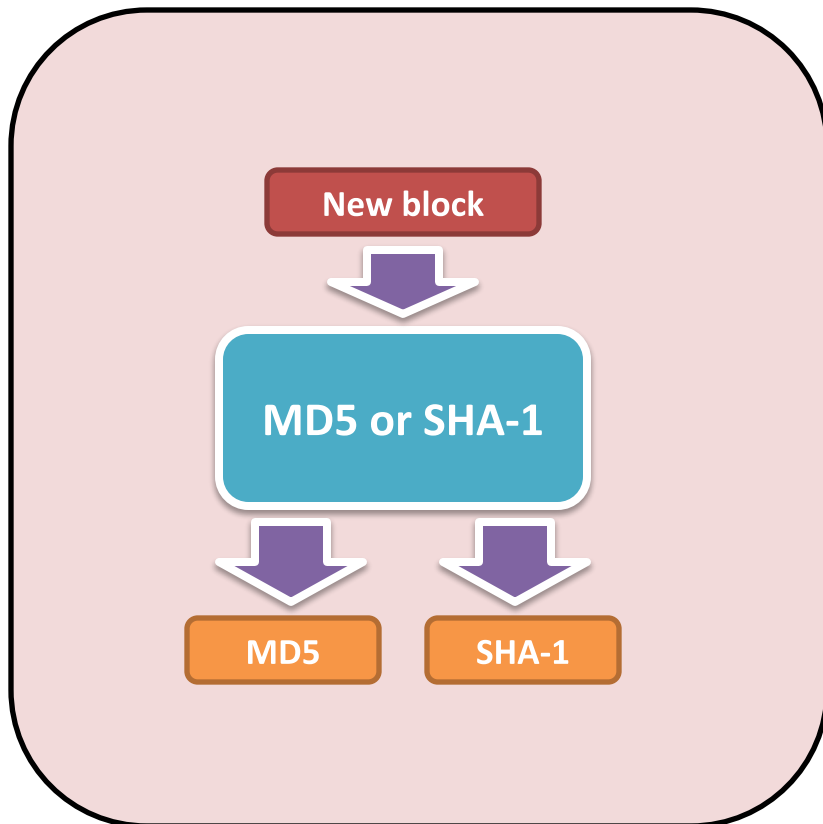
- Deduplication (typically) operates on fixed-size blocks
 - Saves one copy if two fixed-size blocks have the same content
- For VM image storage, deduplication efficiencies similar for fixed-size blocks and variable-size blocks [Jin & Miller, '09]

Basics: Comparing blocks

- How to compare blocks?

Basics: Fingerprints

- How to compare blocks?
- Solution: Use cryptographic hashes (or **fingerprints**)



- Hash-based comparisons
 - Same content → same hash
 - Different content → different hashes with high probability
 - Pros: block comparison reduced to hash comparison
 - Cons: collision may occur, but with negligible probability
- [Quinlan & Dorward, '02]

Basics: How often do I get matches?

Basics: How often do I get matches?

- Let us assume a bad scenario
 - Independent data blocks
 - Data blocks have equal probability to occur

Basics: How often do I get matches?

- Let us assume a bad scenario
 - Independent data blocks
 - Data blocks have equal probability to occur
- Familiar with the birthday paradox?
 - Imagine you have N people and 365 possible birthdays, how many people do you need to have a probability of at least 2 having the same birthdate
 -

Basics: How often do I get matches?

- Let us assume a bad scenario
 - Independent data blocks
 - Data blocks have equal probability to occur
- Familiar with the birthday paradox?
 - Imagine you have N people and 365 possible birthdays, how many people do you need to have a probability of at least 2 having the same birthdate.
 - Number of values to test to have a probability of at least one collision greater than $\frac{1}{2}$ is $N^{1/2}$ for N possible values
 - Block of size $N = 2^m$
 - Need $2^{m/2}$

Basics: How often do I get matches?

- Let us assume a bad scenario
 - Independent data blocks
 - Data blocks have equal probability to occur
- Familiar with the birthday paradox?
 - Imagine you have N people and 365 possible birthdays, how many people do you need to have a probability of at least 2 having the same birthdate.
 - Number of values to test to have a probability of at least one collision greater than $\frac{1}{2}$ is $N^{1/2}$ for N possible values
 - Block of size $N = 2^m$
 - Need $2^{m/2}$

Example: $m = 10$ bits $\rightarrow N = 1024$, need $2^5 = 32$

Basics: How often do I get matches?

- Let us assume a bad scenario
 - Independent data blocks
 - Data blocks have equal probability to occur
- Familiar with the birthday paradox?
 - Imagine you have N people and 365 possible birthdays, how many people do you need to have a probability of at least 2 having the same birthdate.
 - Number of values to test to have a probability of at least one collision greater than $\frac{1}{2}$ is $N^{1/2}$ for N possible values
 - Block of size $N = 2^m$
 - Need $2^{m/2}$

Example: $m = 1000$ bits $\rightarrow N = 2^{1000} \rightarrow 2^{500} \sim 3 \times 10^{150}$

Basics: How often do I get matches?


- Let us assume a bad scenario
 - Independent data blocks
 - Data blocks have equal probability to occur
- Familiar with the birthday paradox?
 - Imagine you have N people and 365 possible birthdays, how many people do you need to have a probability of at least 2 having the same birthdate
 - Number of values to test to have a probability of at least one collision greater than $\frac{1}{2}$ is $N^{1/2}$ for N possible values
 - Block of size $N = 2^m$
 - Need $2^{m/2}$

Typically, there is a lot more correlation in the data :-)

**Basics: How to know if a block
should be deleted?**

Basics: Reference Counts

- How to know if a block should be deleted?
- **Solution:** Keep a **reference count** for each block. Zero means the block is no longer referenced



Inline Deduplication

- How to check if a block being written can be deduplicated with existing blocks?
- Solution: maintain an **index structure**
 - Keep track of fingerprints of existing blocks
- Goal: design of index structure must be efficient in **space** and **speed**
- Two options of keeping an index structure:
 - Putting whole index structure in RAM
 - Putting whole index structure on disk

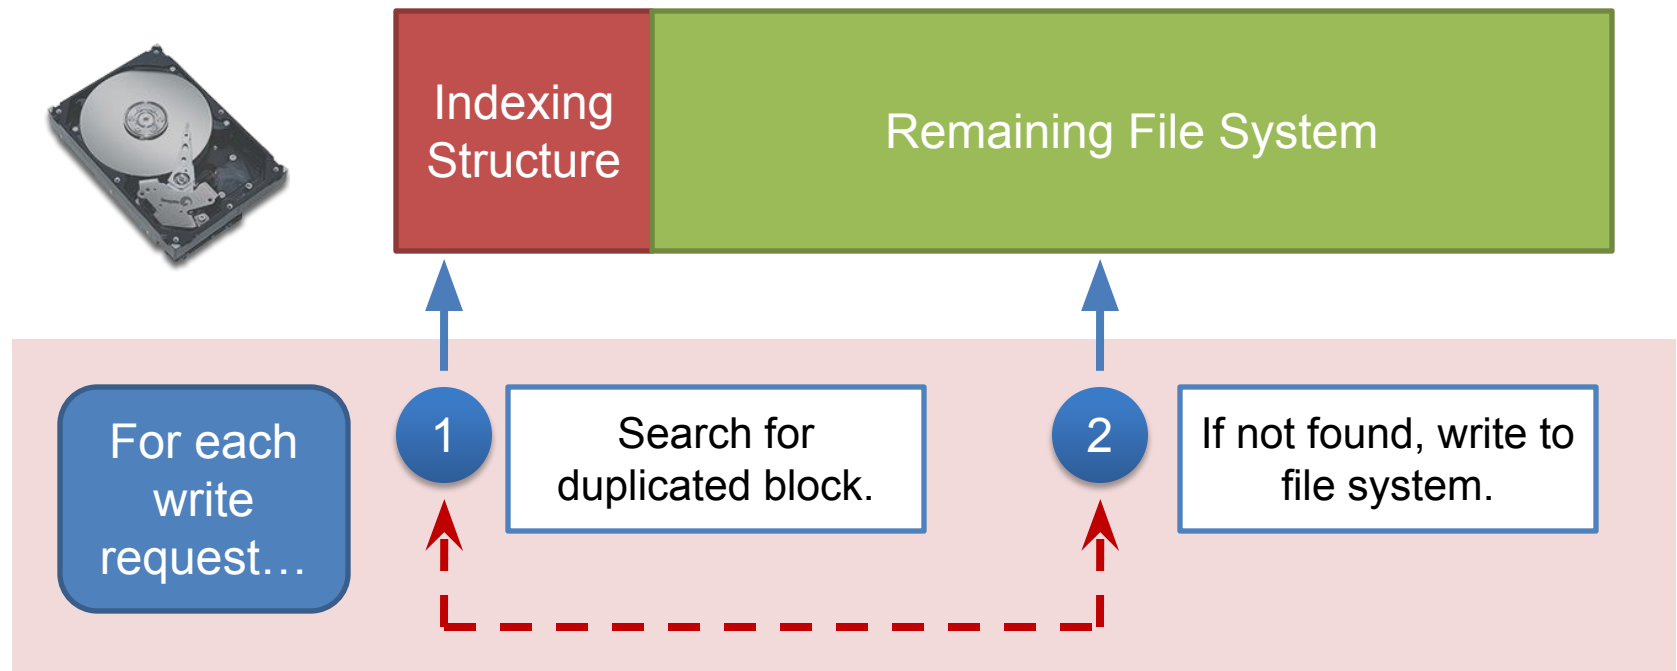
Option 1: Index Structure in RAM

- How about putting whole index structure in **RAM**?
 - Used in existing dedup file systems (e.g., ZFS, OpenDedup)
- Challenge: need large amount of RAM
- Example: per 1TB of disk content

Block Size	4KB
Using MD5 checksum	16 bytes per block
Size of Index	$1\text{TB} / 4\text{KB} \times 16 \text{ bytes} = \mathbf{4GB}$

Option 2: Index Structure on Disk

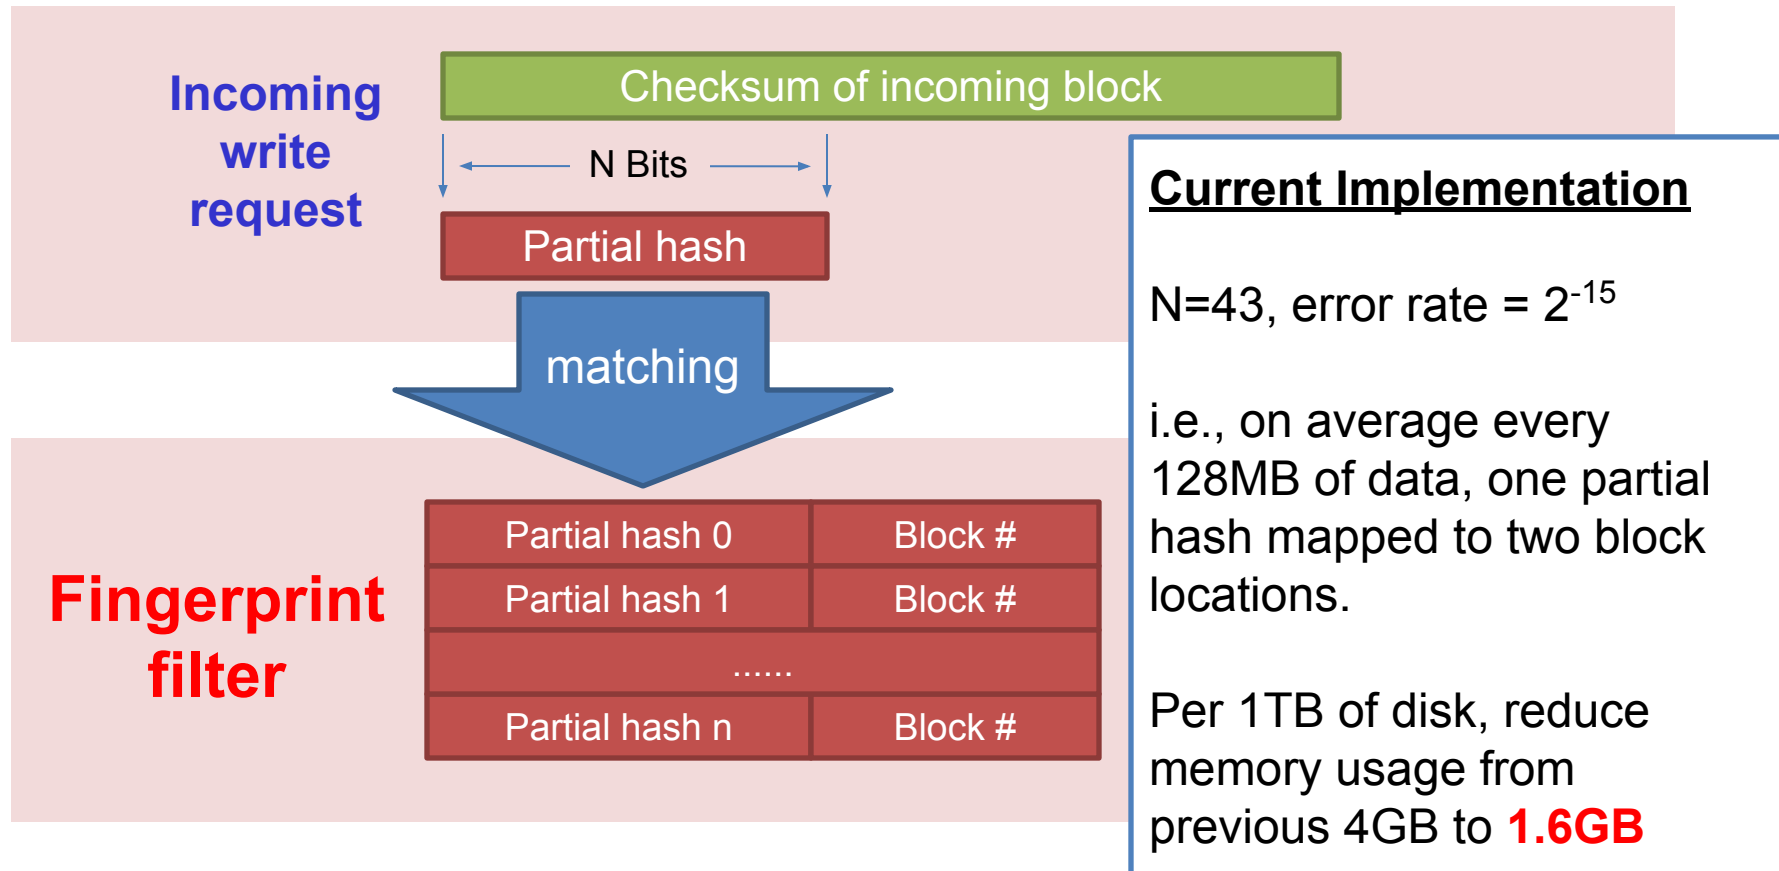
- How about putting whole index structure **on disk**?



- Challenge: updating each data block and its index keeps the disk head moving, which hurts performance


Example: LiveDFS Design

- Store partial fingerprints in memory
 - Infer if same block exists, and where it is “potentially” located



Example: LiveDFS Design

- Store full fingerprints on disk, with **spatial locality**
 - Verify whether the same block is actually located



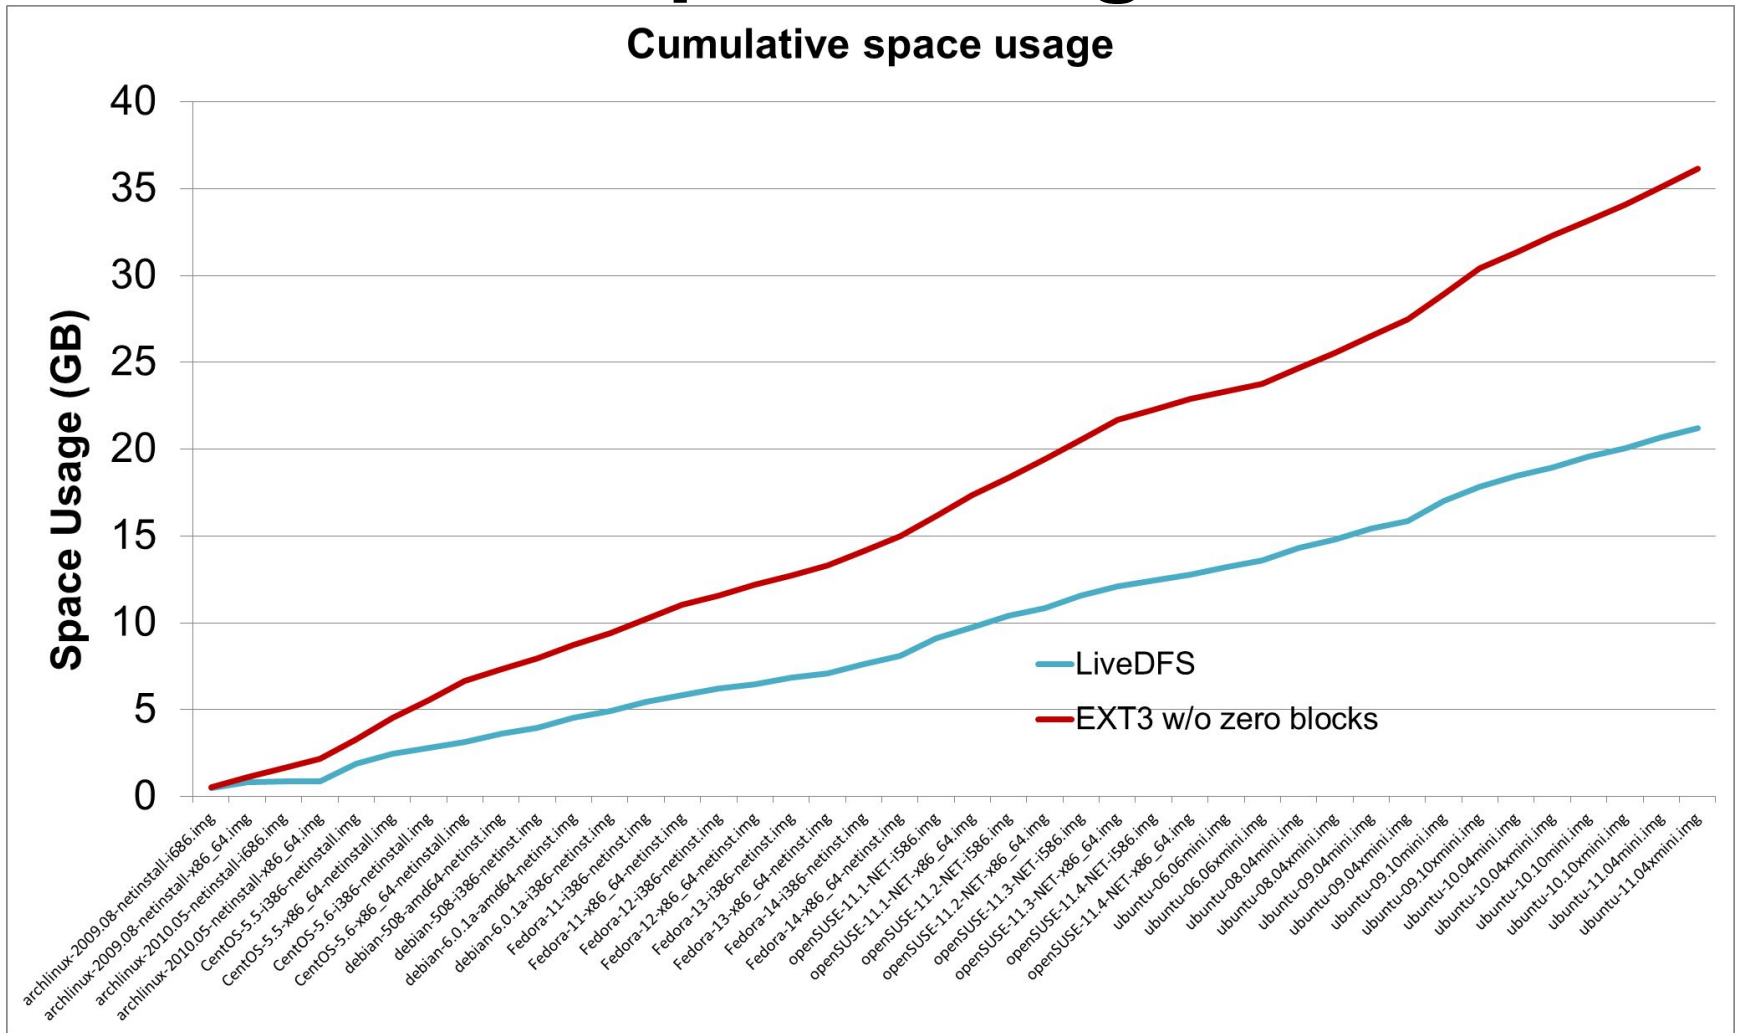
- Partition index structure according to block groups
 - Each block group has a **fingerprint store**
 - Each fingerprint store keeps fingerprints and reference counts for the respective data blocks in the same block group
- Writing with close proximity incurs minimal seeks

Example: LiveDFS Design

- Take-away: LiveDFS arranges fingerprints in memory and on disk according to **underlying file system layout on disk**
- Other features:
 - Prefetching of fingerprint store:
 - load entire fingerprint store of same block group into page cache
 - subsequent writes updates fingerprint store directly in page cache.
 - Journaling:
 - follow Linux file system journaling design
 - enable crash recovery and enhance write performance by combining block writes in batch

Space Usage

Cumulative space usage




LiveDFS saves 40% storage over Ext3


Slide adapted from Patrick Lee

A Generalized View of Deduplication


Generalized Deduplication



Generalized Deduplication

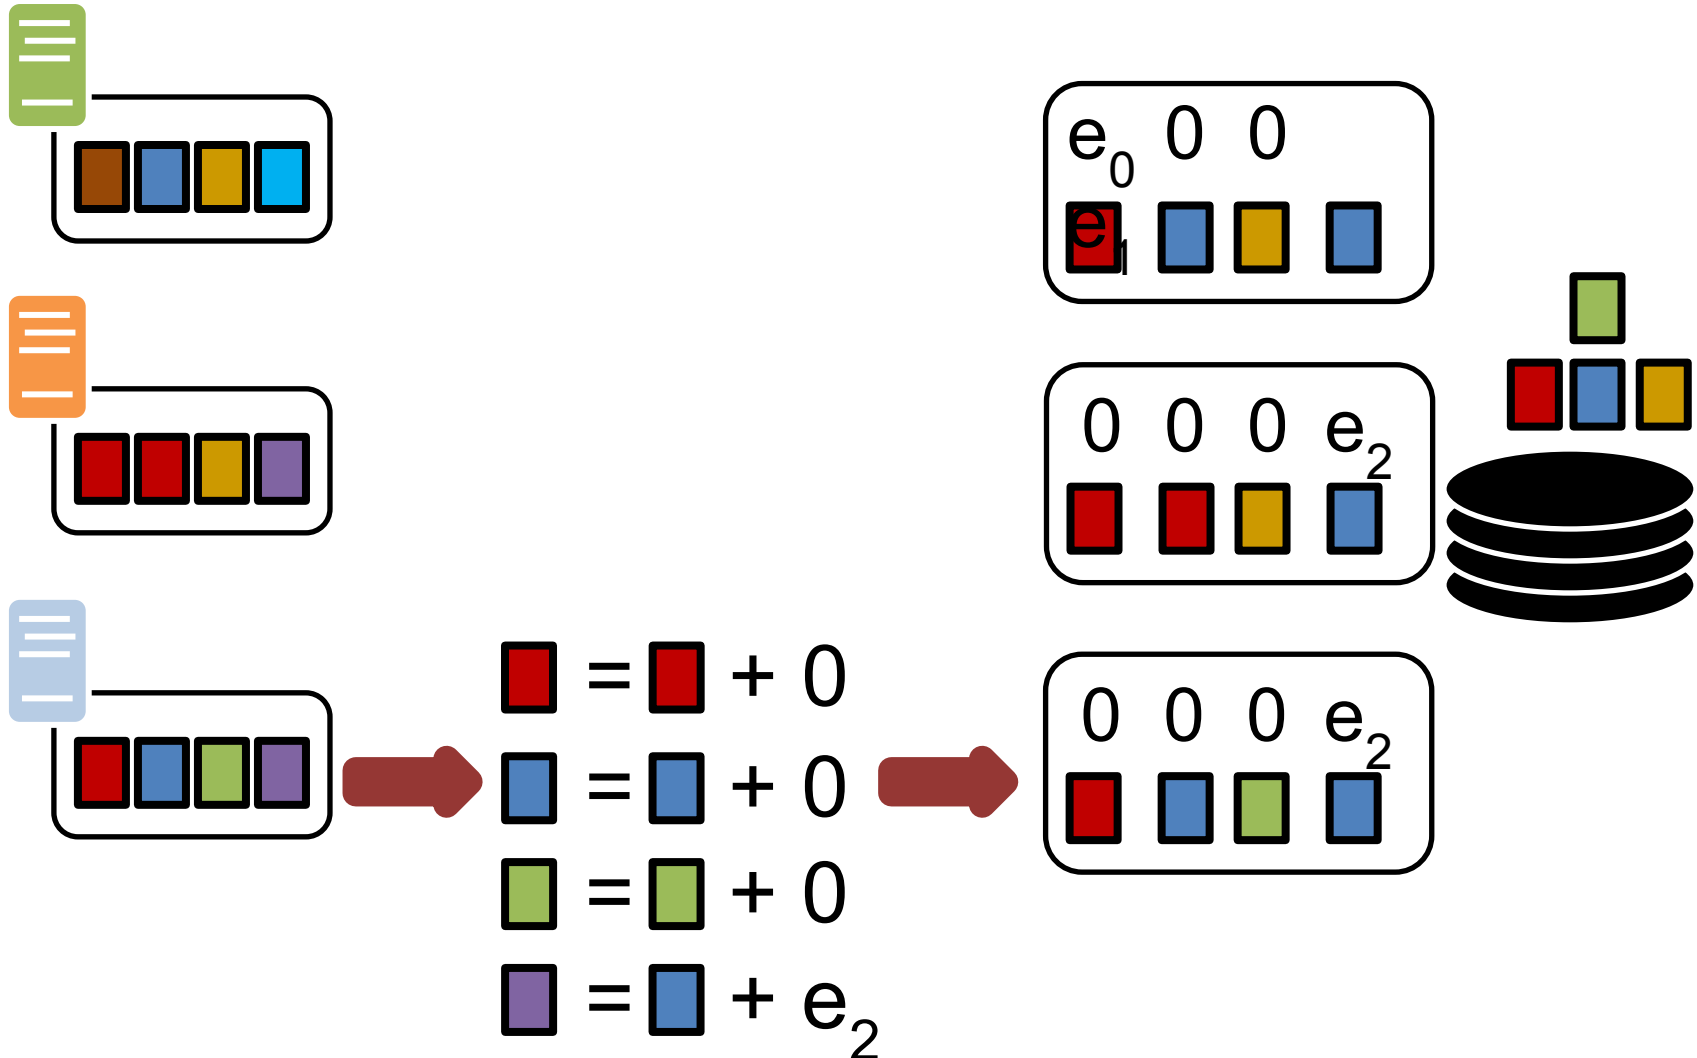


Generalized Deduplication




Generalized Deduplication

Compression gain is $12/4 = 3$




HOW TO DO THE TRANSFORMATION?

- Promotes large amount of matches
- Small deviation (few bits)
- There are a number of options




Example of Transformation

Error Correcting Codes: case where code can correct errors




Example of Transformation

Error Correcting Codes: case where code can correct errors




In our case...

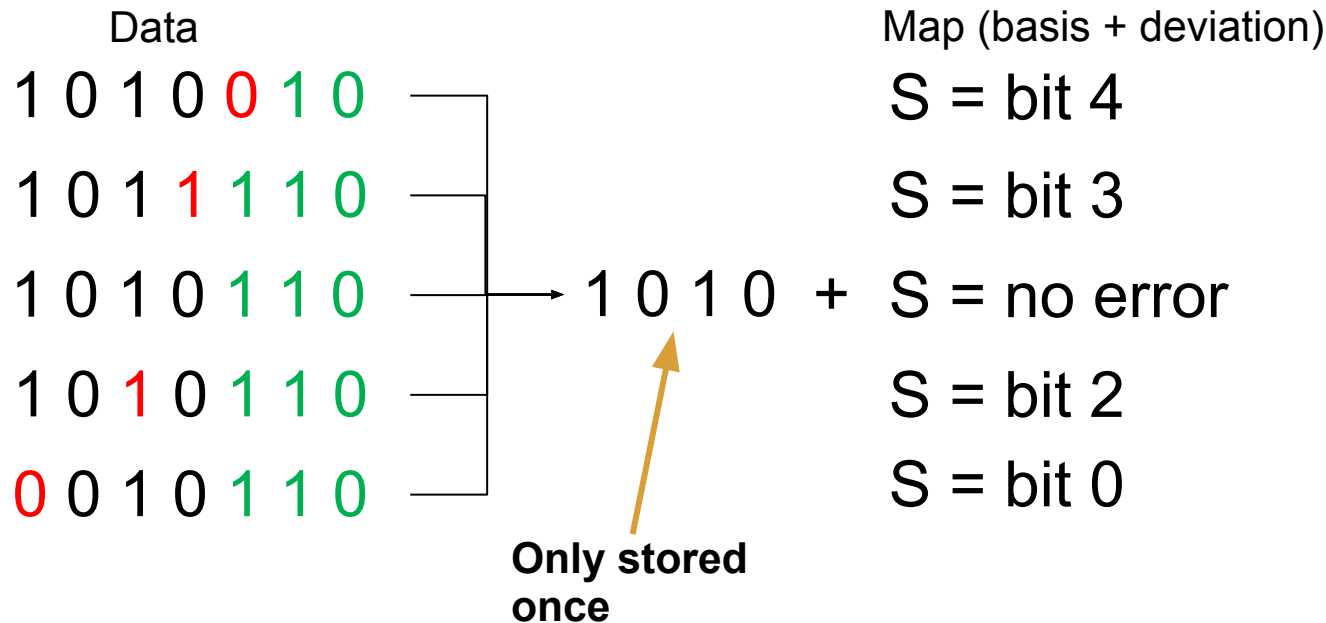


Example of Transformation

Error Correcting Codes: case where code can correct errors




In our case...

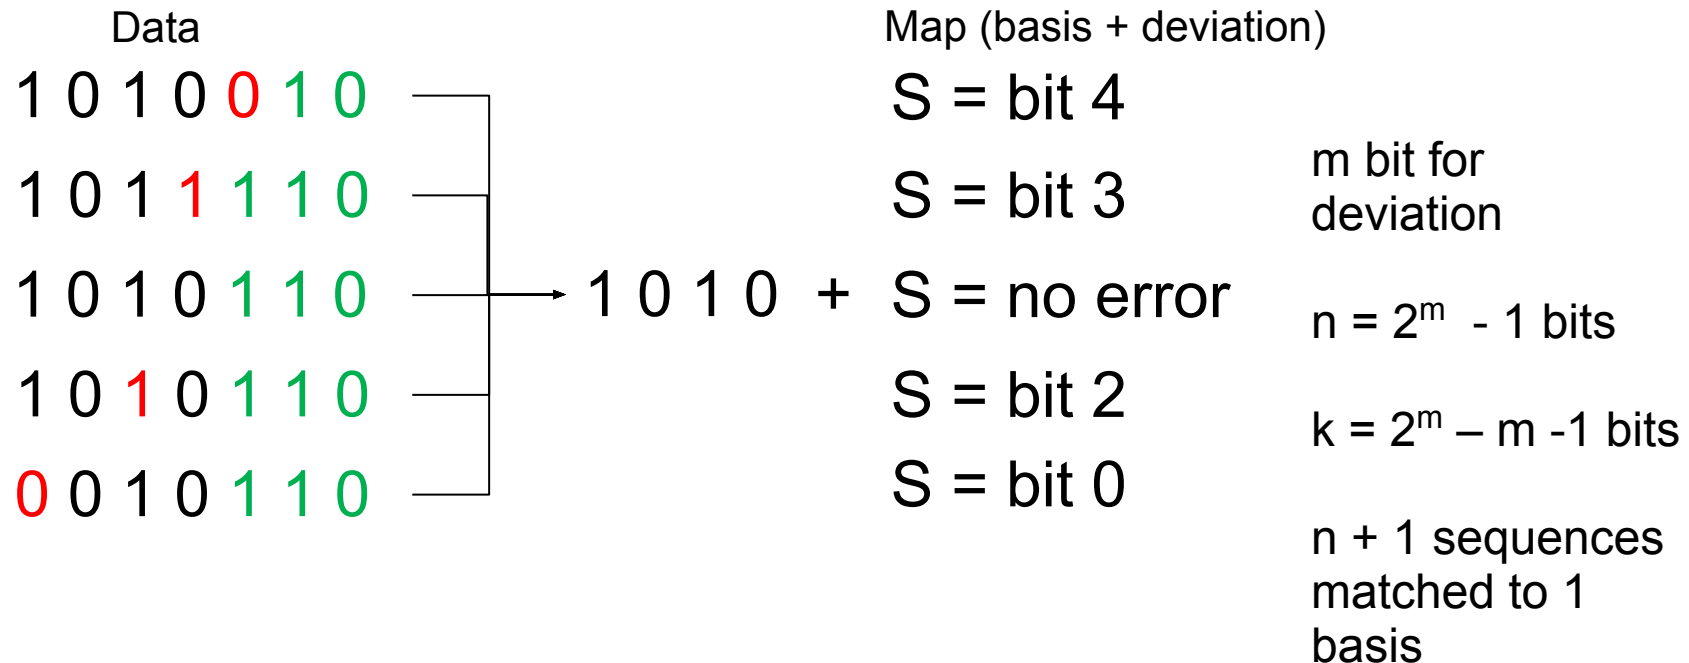


Example of Transformation

Error Correcting Codes: case where code can correct errors




In our case...



HOW TO DO THE TRANSFORMATION?

- Promotes large amount of matches: $n + 1$ matches
- Small deviation (few bits): m bits
- There are a number of options with better characteristics





If $m = 15$ bits


$n \sim 4$ KB

32768 matches


10/10/2016




Compression Starts Earlier




Compression Starts Earlier



Information Theory – proven this speed up



Need Smaller Registry

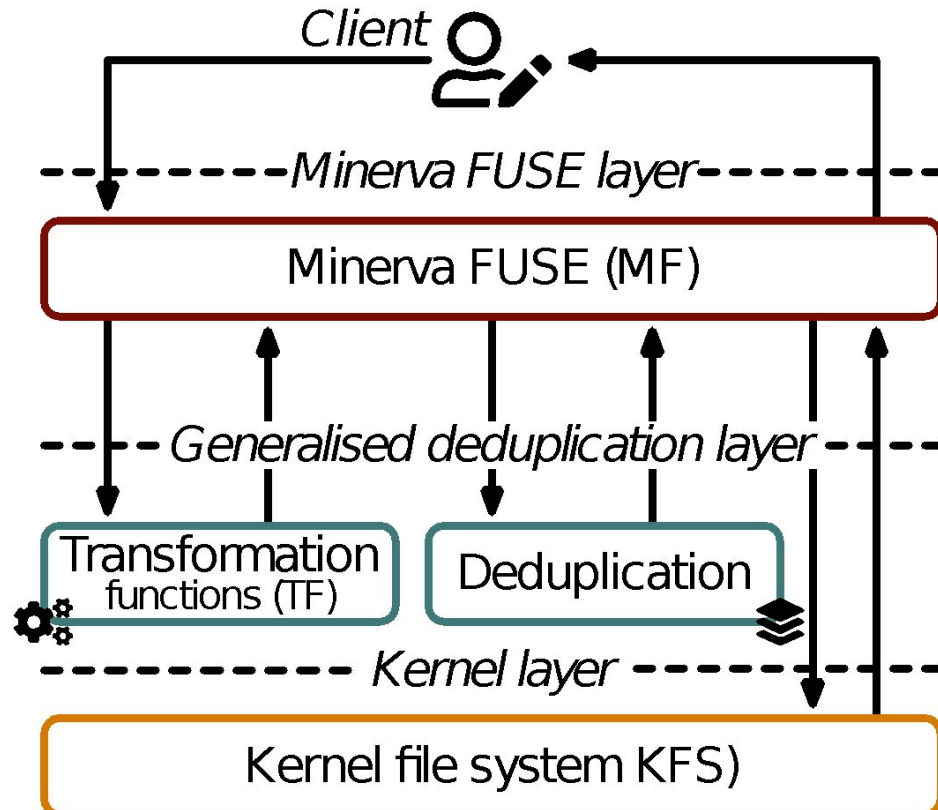


Size of the registry gen dedup vs dedup


ZFS: rule of thumb
5GB of Memory for
every 1TB of data

Can lower this
requirement
(e.g. 10 – 10000
times)

MINERVA-FS



MINERVA-FS



Also... much more
compression
than LiveDFS!

FUSE over EXT4

- ZFS: ~5GB of Memory for every 1TB of data
- MinervaFS: due to GD & implementation is an order of magnitude lower for the same chunk size